

A representative rationale comparison

The same forecasting question and evidence are passed through different rationale parameters

1 Concrete forecasting case

Will China mass-produce humanoid robots by the end of 2025?

Resolves Yes if a China-based company produces >1,000 identical units by 31 Dec 2025.

Retrieved evidence

Robot Report, Sep. 2025: Chinese manufacturers target >1,000 units by the end of 2025.

Other reports describe plans without confirming the threshold.

2 Rationale parameters

Same question, resolution criterion, and evidence

V0 Neutral
Brief rationale

V4 Credibility
Assess source reliability

V8 Temporal
Check event and deadline

Only the requested rationale structure changes

3 Forecasts for this case

No | 0.35
Treats reports as plans

Yes | 0.68
Selects threshold-relevant report

Yes | 0.66
Checks the 2025 deadline

Ground truth: YES | resolved outcome

4 Study-level evaluation

**1,580 questions x 9 prompt conditions
x 3 models x 6 temperatures**

Primary: Brier score, log loss
Diagnostics: accuracy, ECE, validity, rationale quality

Interpretation

The example shows forecast variation under matched prompts. Aggregate metrics test whether structured rationales improve forecasting.

Figure 1: Representative case from the forecasting study (Metaculus 39523; GPT-OSS-120B). The example is illustrative; the paper's main claims are based on paired aggregate evaluation.