

A representative rationale comparison

The same forecasting question and evidence are passed through different rationale parameters

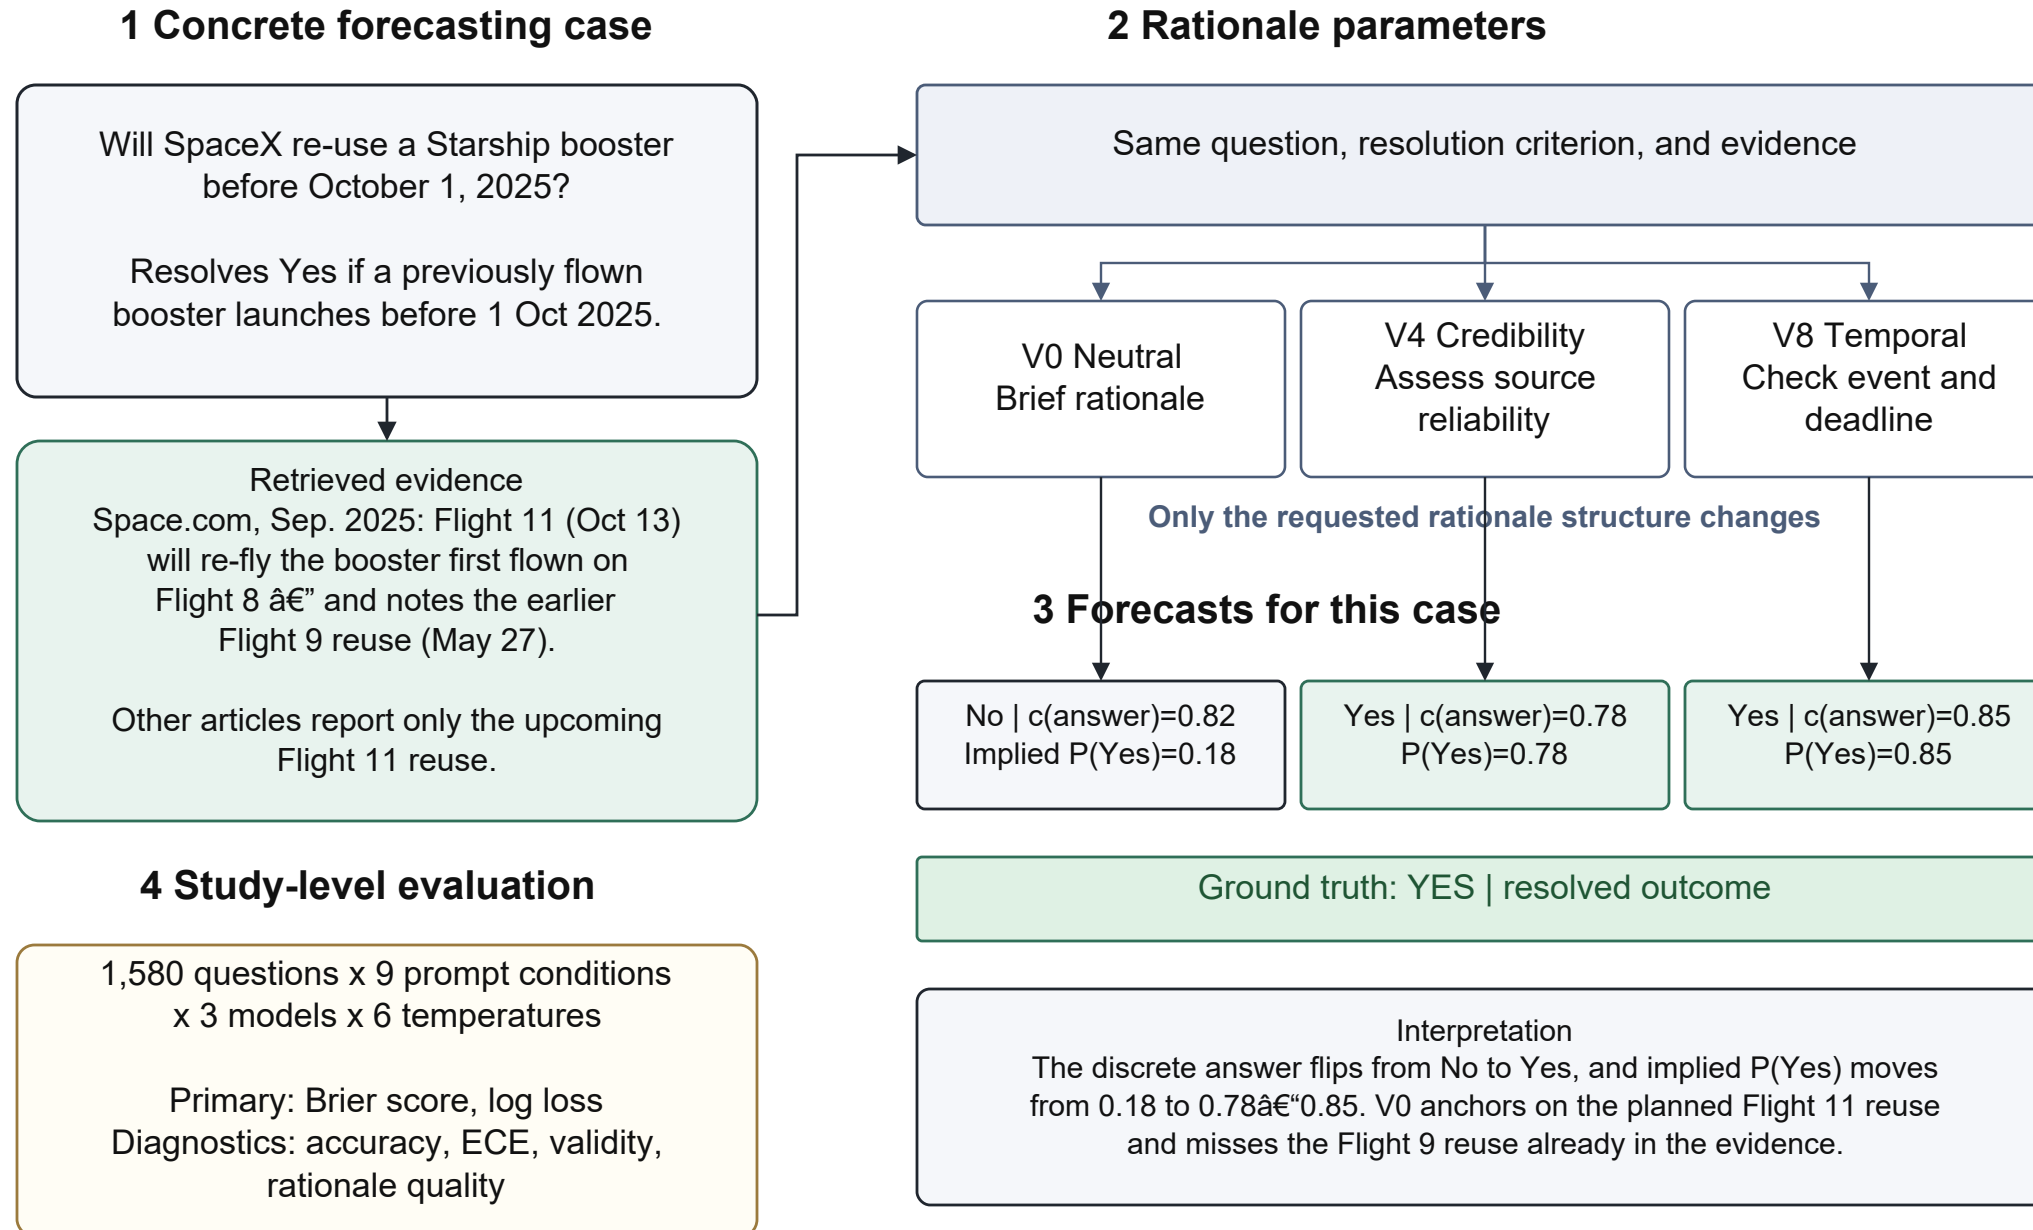


Figure 1: Representative case from the forecasting study (Metaculus 29222; GPT-OSS-120B, T=0.75). Boxes reproduce the raw answer and the confidence in that answer, $c(\text{answer})$. For evaluation, No forecasts are converted using $P(\text{Yes}) = 1 - c(\text{answer})$. Accuracy uses the discrete answer, whereas Brier score and log loss use the implied $P(\text{Yes})$. The example is illustrative; the paper's main claims are based on paired aggregate evaluation.